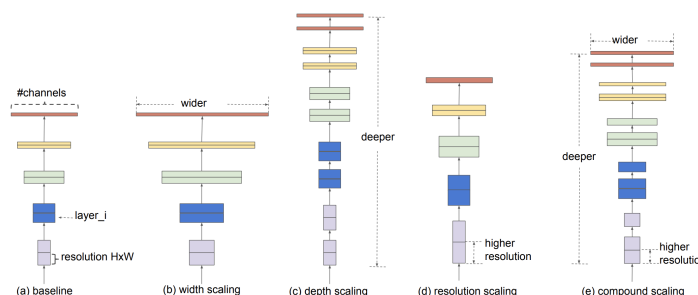




## Question 1

EfficientNet is a convolutional neural network architecture and scaling method that uniformly scales all depth<sup>1</sup>/width<sup>2</sup>/resolution<sup>3</sup> dimensions using a compound coefficient. Unlike conventional practice that arbitrarily scales these factors, the EfficientNet scaling method uniformly scales network width, depth, and resolution with a set of fixed scaling coefficients. For example, we want to use  $2^N$  times more computational resources. In that case, we can increase the network depth by  $\alpha^N$ , width by  $\beta^N$ , and image size by  $\gamma^N$ , where  $\alpha$ ,  $\beta$ , and  $\gamma$  are constant coefficients determined by a small grid search on the original miniature model. EfficientNet uses a compound coefficient PHI to uniformly scales network width, depth, and resolution in a principled way.



The intuition justifies the compound scaling method that if the input image is more extensive, the network needs more layers to increase the receptive field and more channels to capture more fine-grained patterns on the bigger picture.

The base EfficientNet-B0 network is based on the inverted bottleneck residual blocks of MobileNetV2 and squeeze-and-excitation blocks.

EfficientNets also transfer well and achieve state-of-the-art accuracy on CIFAR-100 (91.7%), Flowers (98.8%), and three other transfer learning datasets, with an order of magnitude fewer parameters.


<sup>1</sup>Number of layers (including output but excluding input. E.g., 101). The deeper the network is, the more likely it will experience exploding or vanishing gradients, but it would be more complex and maybe more performant.

<sup>2</sup>Highest number of convolution kernels (channels). As pointed out by Zagoruyko and Komodakis, "wider networks tend to be able to capture more fine-grained features and are easier to train." However, a model too wide and too shallow would have difficulties in capturing higher-level features (e.g., 1024).

<sup>3</sup>Input image's dimension (image height \* image width. e.g. 256 x 256). The higher the resolution, the more likely CNNs will be to capture fine-grained patterns, but the accuracy gain diminishes for very high resolutions (e.g., 560 x 560).

## Question 2

ResNeXt is a homogeneous neural network that reduces the number of hyperparameters required by conventional ResNet. This is achieved by using "cardinality", an additional dimension on top of the width and depth of ResNet. Cardinality defines the size of the set of transformations.

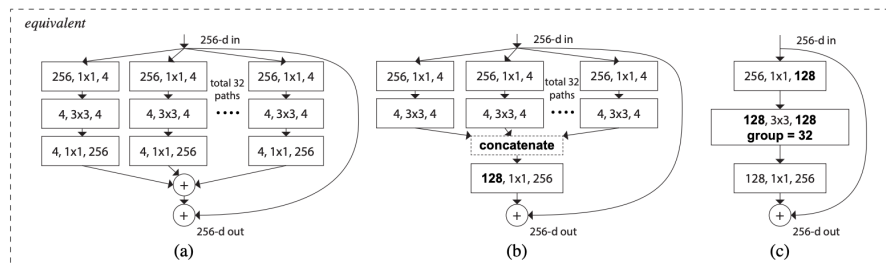


A ResNeXt repeats a building block that aggregates a set of transformations with the same topology. Compared to a ResNet, it exposes a new dimension, cardinality (the size of the set of transformations), as an essential factor in addition to depth and width dimensions.

Two rules define the basic architecture of ResNeXt. First, if the blocks produce same-dimensional spatial maps, they share the same set of hyperparameters, and if at all the spatial map is down-sampled by a factor of 2, the block's width is multiplied by a factor of 2.

| stage     | output  | ResNet-50   | ResNeXt-50 (32×4d)  |
|-----------|---------|---|---|
| conv1     | 112×112 | 7×7, 64, stride 2   | 7×7, 64, stride 2   |
| conv2     | 56×56   | 3×3 max pool, stride 2  | 3×3 max pool, stride 2  |
|           |         | $\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$    | $\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128, C=32 \\ 1 \times 1, 256 \end{bmatrix} \times 3$    |
| conv3     | 28×28   | $\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$  | $\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256, C=32 \\ 1 \times 1, 512 \end{bmatrix} \times 4$    |
| conv4     | 14×14   | $\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$ | $\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512, C=32 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$   |
| conv5     | 7×7     | $\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 1024 \\ 3 \times 3, 1024, C=32 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$ |
|           | 1×1     | global average pool   | global average pool   |
|           |         | 1000-d fc, softmax  | 1000-d fc, softmax  |
| # params. |         | $25.5 \times 10^6$  | $25.0 \times 10^6$  |
| FLOPs     |         | $4.1 \times 10^9$   | $4.2 \times 10^9$   |

Denote the residual block structures, whereas the numbers written adjacent to them refer to the number of stacked blocks. 32 precisely denotes that there are 32 groups in the grouped convolution.



The above network structures explain a grouped convolution and how it trumps the other two network structures.

- denotes a usual ResNeXt block that has already been seen previously. It has a cardinality of 32 and follows the split-transform-merge strategy.
- does seem to be a leaf taken out of Inception-ResNet. However, Inception or Inception-ResNet doesn't have network blocks following the same topology.
- is related to the grouped convolution, which has been proposed in AlexNet architecture.  $32 \times 4$ , as has been seen in (a) and (b), has been replaced with 128 in short, meaning splitting is done by a grouped convolutional layer. Similarly, the transformation is done by the other grouped convolutional layer that does 32 groups of convolutions. Later, concatenation happens.

Inception increases the network space from which the best network is to be chosen via training. Each inception module can capture salient features at different levels. Global attributes are captured by the  $5 \times 5$  Conv layer, while the  $3 \times 3$  Conv layer is prone to capturing distributed features. The max-pooling operation captures low-level features that stand out in a neighborhood. All of these features are extracted and concatenated at a given level before it is fed to the next layer. We leave for the network/training to decide what features hold the most value and weight accordingly. Say if the images in the data set are rich in global features without too many low-level features, then the trained Inception network will have minimal weights corresponding to the  $3 \times 3$  Conv kernel compared to the  $5 \times 5$  Conv kernel.

In the table below, these four CNNs are sorted w.r.t their top-5 accuracy on the Imagenet dataset. The number of trainable parameters and the Floating Point Operations (FLOP) required for a forward pass can also be seen.

As a last point, while Inception focuses on computational cost, ResNeXt focuses on computational accuracy. Intuitively, deeper networks should not perform worse than shallower networks. Still, in practice, the deeper networks performed worse than the more superficial ones, caused by overfitting and an optimization problem.

| Comparison |      |                          |               |            |       |
|------------|------|--------------------------|---------------|------------|-------|
| Network    | Year | Salient Feature          | top5 accuracy | Parameters | FLOP  |
| AlexNet    | 2012 | Deeper                   | 84.70%        | 62M        | 1.5B  |
| VGGNet     | 2014 | Fixed-size kernels       | 92.30%        | 138M       | 19.6B |
| Inception  | 2014 | Wider - Parallel kernels | 93.30%        | 6.4M       | 2B    |
| ResNet-152 | 2015 | Shortcut connections     | 95.51%        | 60.3M      | 11B   |

### Question 3

# 1 Natural scene classification

```
[1]: import tensorflow as tf
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from google.colab import output
print(tf.version.VERSION)
```

2.9.2

```
[2]: print('Default GPU Device: {}'.format(tf.test.gpu_device_name()))
```

Default GPU Device: /device:GPU:0

## 2 Setting up Random Seed

```
[3]: RANDOM_SEED: int = 42
```

```
[4]: import random
import os
os.environ['PYTHONHASHSEED'] = str(RANDOM_SEED)
random.seed(RANDOM_SEED)
tf.random.set_seed(RANDOM_SEED)
tf.experimental.numpy.random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)
```

## 3 Get class names (labels)

```
[5]: !kaggle datasets download -d puneet6060/intel-image-classification
!unzip intel-image-classification.zip
output.clear()
```

```
[6]: import pathlib
data_dir = pathlib.Path("/content/seg_train/seg_train/")
class_names = np.array(sorted([item.name for item in data_dir.glob("*")])).
    →tolist() # created a list of class_names from the subdirector
class_names
```

```
[6]: ['buildings', 'forest', 'glacier', 'mountain', 'sea', 'street']
```

## 4 Visualize random image from dataset

```
[7]: # visualize random image from train set
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

target_class = random.choice(class_names)
target_folder: str = "/content/seg_train/seg_train/" + target_class

# Get a random image path
random_image = random.sample(os.listdir(target_folder), 1)

# Read in the image and plot it using matplotlib
img = mpimg.imread(target_folder + "/" + random_image[0])
plt.imshow(img)
plt.title(f"{target_class}: {random_image[0]}")

print(f"Image shape: {img.shape}")
```

Image shape: (150, 150, 3)



## 5 Setting up Hyperparameters

```
[8]: BATCH_SIZE: int = 32
      EPOCHS: int = 12
      IMAGE_SIZE = (150, 150)
      AUGMENTATION_FACTOR: float = 0.2
      LABEL_MODEL: str = "categorical"

      TRAIN_DIR: str = "/content/seg_train/seg_train"
      TEST_DIR: str = "/content/seg_test/seg_test"
```

## 6 Preprocessing data

```
[9]: print("Training data :")
      train_data = tf.keras.preprocessing.image_dataset_from_directory(
          directory=TRAIN_DIR,
          label_mode=LABEL_MODEL,
          image_size=IMAGE_SIZE,
          batch_size=BATCH_SIZE,
          seed=RANDOM_SEED,
          shuffle=True
      )

      print("Testing data :")
      test_data = tf.keras.preprocessing.image_dataset_from_directory(
          directory=TEST_DIR,
          label_mode=LABEL_MODEL,
          image_size=IMAGE_SIZE,
          batch_size=BATCH_SIZE,
          seed=RANDOM_SEED,
          shuffle=False
      )
```

```
Training data :
Found 14034 files belonging to 6 classes.
Testing data :
Found 3000 files belonging to 6 classes.
```

## 7 Visualize augmentation layer

```
[10]: from tensorflow.keras.layers import Dense, RandomFlip, RandomRotation, \
      ↪RandomZoom, RandomWidth, RandomHeight, Rescaling
      from tensorflow.keras import Sequential
      from tensorflow.keras.activations import softmax
```

```
[11]: augmentation_layer = Sequential([
        RandomFlip("horizontal"),
        RandomRotation(AUGMENTATION_FACTOR),
        RandomZoom(AUGMENTATION_FACTOR),
        RandomWidth(AUGMENTATION_FACTOR),
        RandomHeight(AUGMENTATION_FACTOR),
        Rescaling(1./255)
    ], name="augmentation_layer")

augmentation_layer
```

```
[11]: <keras.engine.sequential.Sequential at 0x7f61dc305e80>
```

## 8 Visualize augmentation layer

```
[12]: target_dir = TRAIN_DIR
target_class = random.choice(class_names)
target_dir = f"{target_dir}/{target_class}"
random_image = random.choice(os.listdir(target_dir))
random_image_path = target_dir + "/" + random_image

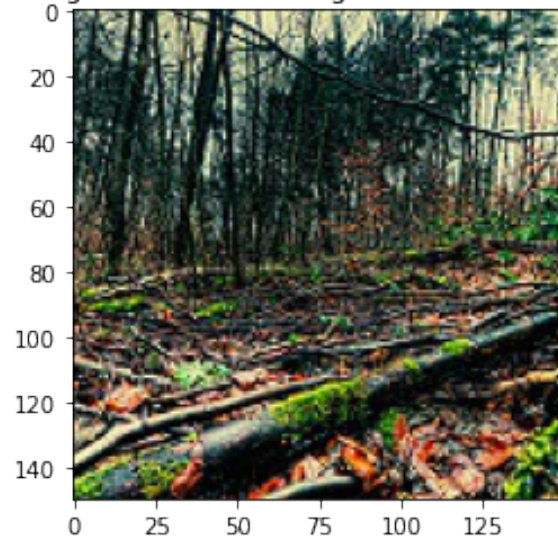
print(random_image_path)
# Read in the random image
img = mpimg.imread(random_image_path)
plt.title(f"Original random image from class: {target_class}")
# plt.axis(False)
plt.imshow(img);

# Now lets plot our augmented random image
augmented_image = augmentation_layer(img, training=True)
plt.figure()
plt.title(f"augmented random image from class: {target_class}")
plt.imshow(augmented_image)
```

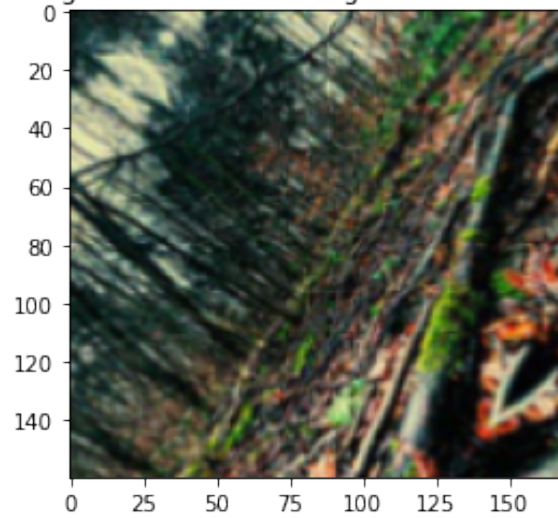
```
/content/seg_train/seg_train/forest/2923.jpg
```

```
[12]: <matplotlib.image.AxesImage at 0x7f61dc244730>
```

Original random image from class: forest



augmented random image from class: forest





## 9 Creating the backbone base model using ResNet50V2

```
[13]: # Turn off all warnings except for errors
tf.get_logger().setLevel('ERROR')

# set seed
tf.random.set_seed(RANDOM_SEED)

base_model = tf.keras.applications.ResNet50V2(include_top=False)
base_model.trainable = False

# Creating input layer
input_layer = tf.keras.layers.Input(shape=IMAGE_SIZE + (3,), name="input_layer")
x = augmentation_layer(input_layer)
x = base_model(x, training=False)
x = tf.keras.layers.
    ↳GlobalAveragePooling2D(name="global_average_pooling_2d_layer")(x)
output_layer = Dense(len(class_names), activation=softmax,
    ↳name="output_layer")(x)

model_1 = tf.keras.Model(input_layer, output_layer)
```

```
[14]: model_1.summary()
```

Model: "model"

| Layer (type)   | Output Shape             | Param #  |
|--|--------------------------|----------|
| input_layer (InputLayer)                                     | [(None, 150, 150, 3)]    | 0        |
| augmentation_layer (Sequential)                              | (None, None, 3)          | 0        |
| resnet50v2 (Functional)                                      | (None, None, None, 2048) | 23564800 |
| global_average_pooling_2d_1<br>ayer (GlobalAveragePooling2D) | (None, 2048)             | 0        |
| output_layer (Dense)   | (None, 6)                | 12294    |

=====  
Total params: 23,577,094  
Trainable params: 12,294  
Non-trainable params: 23,564,800  
=====

## 10 Tensorflow Callbacks

```
[15]: early_stopping_callback = tf.keras.callbacks.EarlyStopping(monitor="val_loss",
                                                                patience=3)
reduce_lr_callback = tf.keras.callbacks.ReduceLROnPlateau(monitor="val_loss",
                                                         factor=0.2,
                                                         patience=2,
                                                         verbose=1,
                                                         min_lr=1e-7)

# set checkpoint path
checkpoint_path = "checkpoint_weights/checkpoint.cpk"

# Create a ModelCheckpoint callback that saves the model's weights only
checkpoint_callback = tf.keras.callbacks.ModelCheckpoint(
    filepath=checkpoint_path,
    save_weights_only=True,
    save_best_only=True,
    save_freq="epoch", #save every epoch
    verbose=1
)
```

## 11 Compile the model

```
[16]: model_1.compile(
    loss=tf.keras.losses.CategoricalCrossentropy(),
    optimizer=tf.keras.optimizers.Adam(),
    metrics=["accuracy"]
)
```

## 12 Fit the model

```
[17]: history_1 = model_1.fit(
    train_data,
    epochs=EPOCHS,
    steps_per_epoch=len(train_data),
    validation_data=test_data,
    validation_steps=int(0.15 * len(test_data)),
    callbacks=[
        early_stopping_callback,
        reduce_lr_callback,
        checkpoint_callback
    ]
)
```

Epoch 1/12

439/439 [=====] - ETA: 0s - loss: 0.5573 - accuracy: 0.7953

Epoch 1: val\_loss improved from inf to 0.36231, saving model to checkpoint\_weights/checkpoint.cpk

439/439 [=====] - 100s 214ms/step - loss: 0.5573 - accuracy: 0.7953 - val\_loss: 0.3623 - val\_accuracy: 0.8661 - lr: 0.0010

Epoch 2/12

439/439 [=====] - ETA: 0s - loss: 0.4293 - accuracy: 0.8448

Epoch 2: val\_loss improved from 0.36231 to 0.18632, saving model to checkpoint\_weights/checkpoint.cpk

439/439 [=====] - 54s 124ms/step - loss: 0.4293 - accuracy: 0.8448 - val\_loss: 0.1863 - val\_accuracy: 0.9241 - lr: 0.0010

Epoch 3/12

439/439 [=====] - ETA: 0s - loss: 0.4066 - accuracy: 0.8539

Epoch 3: val\_loss did not improve from 0.18632

439/439 [=====] - 51s 115ms/step - loss: 0.4066 - accuracy: 0.8539 - val\_loss: 0.3115 - val\_accuracy: 0.8906 - lr: 0.0010

Epoch 4/12

439/439 [=====] - ETA: 0s - loss: 0.3944 - accuracy: 0.8555

Epoch 4: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 4: val\_loss did not improve from 0.18632

439/439 [=====] - 48s 109ms/step - loss: 0.3944 - accuracy: 0.8555 - val\_loss: 0.2292 - val\_accuracy: 0.9085 - lr: 0.0010

Epoch 5/12

439/439 [=====] - ETA: 0s - loss: 0.3550 - accuracy: 0.8724

Epoch 5: val\_loss did not improve from 0.18632

439/439 [=====] - 46s 104ms/step - loss: 0.3550 - accuracy: 0.8724 - val\_loss: 0.2607 - val\_accuracy: 0.8996 - lr: 2.0000e-04

```
[18]: # load in saved model wights and evaluate model
      model_1.load_weights(checkpoint_path)
```

```
[18]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at
      0x7f61c2432fa0>
```

```
[19]: model_1.evaluate(test_data)
```

94/94 [=====] - 6s 62ms/step - loss: 0.3210 - accuracy: 0.8810

```
[19]: [0.3210234045982361, 0.8809999823570251]
```

```
[20]: import matplotlib.pyplot as plt
def plot_loss_curve(history):
    loss = history_1.history['loss']
    val_loss = history_1.history['val_loss']


    accuracy = history_1.history["accuracy"]
    val_accuracy = history_1.history["val_accuracy"]


    epochs = range(len(history_1.history['loss']))

    # Plot loss
    plt.plot(epochs, loss, label='training_loss')
    plt.plot(epochs, val_loss, label='val_loss')
    plt.title('Loss')
    plt.xlabel('Epochs')
    plt.legend()

    # Plot accuracy
    plt.figure()
    plt.plot(epochs, accuracy, label='training_accuracy')
    plt.plot(epochs, val_accuracy, label='val_accuracy')
    plt.title('Accuracy')
    plt.xlabel('Epochs')
    plt.legend()
```

```
[21]: plot_loss_curve(history_1)
```





```
[22]: # Turn off all warnings except for errors
```

```
tf.get_logger().setLevel('ERROR')
```

```
# set seed
```

```
tf.random.set_seed(RANDOM_SEED)
```

```
base_model.trainable = True
```

```
for layer in base_model.layers[:-20]:
```

```
    layer.trainable = False
```

```
[23]: model_1.compile(
```

```
    loss=tf.keras.losses.CategoricalCrossentropy(),
```

```
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
```

```
    metrics=["accuracy"]
```

```
)
```

```
[24]: history_1.epoch[-1]
```

```
[24]: 4
```

[25]: *# fine tune for another 5 epochs*

```
fine_tune_epochs = EPOCHS + 10

# Refit the model (same as model_2 except with more trainable layers)
history_2 = model_1.fit(
    train_data,
    epochs=fine_tune_epochs,
    validation_data=test_data,
    steps_per_epoch=len(train_data),
    validation_steps=int(0.15 * len(test_data)),
    initial_epoch=history_1.epoch[-1],
    callbacks=[
        early_stopping_callback,
        reduce_lr_callback,
        checkpoint_callback
    ]
)
```

Epoch 5/22

439/439 [=====] - ETA: 0s - loss: 0.3821 - accuracy: 0.8612

Epoch 5: val\_loss improved from 0.18632 to 0.16137, saving model to checkpoint\_weights/checkpoint.cpk

439/439 [=====] - 58s 122ms/step - loss: 0.3821 - accuracy: 0.8612 - val\_loss: 0.1614 - val\_accuracy: 0.9397 - lr: 1.0000e-04

Epoch 6/22

439/439 [=====] - ETA: 0s - loss: 0.3087 - accuracy: 0.8887

Epoch 6: val\_loss did not improve from 0.16137

439/439 [=====] - 49s 112ms/step - loss: 0.3087 - accuracy: 0.8887 - val\_loss: 0.2182 - val\_accuracy: 0.9062 - lr: 1.0000e-04

Epoch 7/22

439/439 [=====] - ETA: 0s - loss: 0.2792 - accuracy: 0.9009

Epoch 7: ReduceLROnPlateau reducing learning rate to 1.9999999494757503e-05.

Epoch 7: val\_loss did not improve from 0.16137

439/439 [=====] - 48s 108ms/step - loss: 0.2792 - accuracy: 0.9009 - val\_loss: 0.2073 - val\_accuracy: 0.9196 - lr: 1.0000e-04

Epoch 8/22

439/439 [=====] - ETA: 0s - loss: 0.2224 - accuracy: 0.9195

Epoch 8: val\_loss did not improve from 0.16137

439/439 [=====] - 48s 109ms/step - loss: 0.2224 - accuracy: 0.9195 - val\_loss: 0.2202 - val\_accuracy: 0.9152 - lr: 2.0000e-05

```
[26]: # Let's create a function to compare training histories
def compare_historys(original_history, new_history, initial_epochs=5, metric:
    str = "accuracy"):
    """
    Compares two TensorFlow History objects.
    """
    # Get original history measurements
    acc = original_history.history[metric]
    loss = original_history.history["loss"]

    val_acc = original_history.history[f"val_{metric}"]
    val_loss = original_history.history["val_loss"]



    # Combine original history metrics with new_history metrics
    total_acc = acc + new_history.history[metric]
    total_loss = loss + new_history.history["loss"]

    total_val_acc = val_acc + new_history.history[f"val_{metric}"]
    total_val_loss = val_loss + new_history.history["val_loss"]

    # Make plot for accuracy
    plt.figure(figsize=(8, 8))
    plt.subplot(2, 1, 1)
    plt.plot(total_acc, label="Training Accuracy")
    plt.plot(total_val_acc, label="Val Accuracy")
    plt.plot([initial_epochs-1, initial_epochs-1], plt.ylim(), label="Start Fine_
    Tuning")
    plt.legend(loc="lower right")
    plt.title("Training and Validation Accuracy")

    # Make plot for loss
    plt.figure(figsize=(8, 8))
    plt.subplot(2, 1, 2)
    plt.plot(total_loss, label="Training Loss")
    plt.plot(total_val_loss, label="Val Loss")
    plt.plot([initial_epochs-1, initial_epochs-1], plt.ylim(), label="Start Fine_
    Tuning")
    plt.legend(loc="upper right")
    plt.title("Training and Validation Loss")
```

```
[27]: compare_historys(history_1, history_2)
```



```
[28]: # load in saved model weights and evaluate model
model_1.load_weights(checkpoint_path)
```

```
[28]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at
0x7f61dc04cfa0>
```

```
[29]: model_1.evaluate(test_data)
```

```
94/94 [=====] - 5s 54ms/step - loss: 0.2610 - accuracy:
0.9000
```

```
[29]: [0.261038601398468, 0.8999999761581421]
```

## 13 Making predictions

```
[30]: # Make predictions with model
preds_probs = model_1.predict(test_data)
preds_probs[:10]
```

```
94/94 [=====] - 5s 50ms/step
```



```
[30]: array([[9.9990058e-01, 4.2831164e-08, 2.7188033e-05, 9.7682114e-06,
          1.0723115e-05, 5.1790419e-05],
          [9.8721343e-01, 1.4112474e-05, 9.3190138e-05, 8.5504325e-05,
          2.6390061e-04, 1.2329934e-02],
          [9.9800485e-01, 3.1354713e-07, 9.3854014e-06, 9.0478679e-05,
          7.1727110e-05, 1.8232594e-03],
          [9.9022043e-01, 1.3345740e-06, 7.9958641e-05, 1.2207507e-03,
          4.3091490e-03, 4.1683889e-03],
          [9.8196453e-01, 4.2899876e-05, 2.2451412e-03, 2.0842291e-03,
          1.2903975e-02, 7.5932709e-04],
          [3.9960101e-01, 4.8900376e-05, 2.2721787e-03, 1.4422562e-02,
          5.2975065e-01, 5.3904641e-02],
          [9.9230701e-01, 1.5337193e-05, 2.9103007e-04, 7.2559790e-04,
          7.1279053e-04, 5.9482004e-03],
          [2.4947867e-01, 8.3858497e-05, 6.9935521e-04, 8.5182366e-04,
          2.8611184e-03, 7.4602515e-01],
          [9.9140406e-01, 4.2123756e-06, 3.9994295e-04, 2.8560358e-05,
          6.1300234e-04, 7.5500770e-03],
          [9.9613476e-01, 5.0047133e-06, 5.2291594e-05, 1.3840073e-05,
          3.2360532e-04, 3.4704928e-03]], dtype=float32)
```

```
[31]: # Get the pred classes of each label
pred_classes = preds_probs.argmax(axis=1)
pred_classes[:10]
```

```
[31]: array([0, 0, 0, 0, 0, 4, 0, 5, 0, 0])
```

```
[32]: # To get our test labels we need to unravel our test_data BatchDataset
y_labels = []
for images, labels in test_data.unbatch():
    y_labels.append(labels.numpy().argmax()) # currently test labels look like: ↪
    → [0, 0, 0, 1, .... 0, 0], we want the index value where the "1" occurs
y_labels[:10] # look at the first 10
```

```
[32]: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

## 14 Reports

```
[33]: from sklearn.metrics import accuracy_score
sklearn_accuracy = accuracy_score(y_true=y_labels,
                                   y_pred=pred_classes)
sklearn_accuracy
```

```
[33]: 0.9
```

```
[34]: import itertools
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import confusion_matrix

# We need to make some changes to our make_confusion_matrix function to ensure
→the x-labels print vertically
def make_confusion_matrix(y_true, y_pred, classes=None, figsize=(10, 10),
→text_size=15, norm=False, savefig=False):
    """Makes a labelled confusion matrix comparing predictions and ground truth
    →labels.

    If classes is passed, confusion matrix will be labelled, if not, integer class
    →values
    will be used.

    Args:
        y_true: Array of truth labels (must be same shape as y_pred).
        y_pred: Array of predicted labels (must be same shape as y_true).
        classes: Array of class labels (e.g. string form). If `None`, integer labels
    →are used.
        figsize: Size of output figure (default=(10, 10)).
        text_size: Size of output figure text (default=15).
        norm: normalize values or not (default=False).
        savefig: save confusion matrix to file (default=False).

    Returns:
        A labelled confusion matrix plot comparing y_true and y_pred.


    Example usage:
        make_confusion_matrix(y_true=test_labels, # ground truth test labels
                              y_pred=y_preds, # predicted labels
                              classes=class_names, # array of class label names
                              figsize=(15, 15),
                              text_size=10)

    """
    # Create the confusion matrix
    cm = confusion_matrix(y_true, y_pred)
    cm_norm = cm.astype("float") / cm.sum(axis=1)[:, np.newaxis] # normalize it
    n_classes = cm.shape[0] # find the number of classes we're dealing with

    # Plot the figure and make it pretty
    fig, ax = plt.subplots(figsize=figsize)
    cax = ax.matshow(cm, cmap=plt.cm.Blues) # colors will represent how 'correct'
    →a class is, darker == better
    fig.colorbar(cax)

    # Are there a list of classes?
    if classes:
        labels = classes
    else:
        labels = np.arange(cm.shape[0])
```

```
[35]: make_confusion_matrix(y_true=y_labels,
                             y_pred=pred_classes,
                             classes=class_names,
                             figsize=(100, 100),
                             text_size=50)
```



```
[36]: from sklearn.metrics import classification_report
print(classification_report(y_true=y_labels,
                             y_pred=pred_classes))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.87      | 0.94   | 0.91     | 437     |
| 1            | 0.99      | 0.99   | 0.99     | 474     |
| 2            | 0.83      | 0.86   | 0.85     | 553     |
| 3            | 0.87      | 0.82   | 0.84     | 525     |
| 4            | 0.91      | 0.92   | 0.91     | 510     |
| 5            | 0.94      | 0.89   | 0.91     | 501     |
| accuracy     |           |        | 0.90     | 3000    |
| macro avg    | 0.90      | 0.90   | 0.90     | 3000    |
| weighted avg | 0.90      | 0.90   | 0.90     | 3000    |

```
[37]: from sklearn.metrics import classification_report
import pandas as pd

def get_f1_score_on_every_class_name(y_labels, y_true, class_names):
    """Return f1 score on every class name as a dataframe

    Args:
        y_labels (_type_): y_true of test_
        y_pred (_type_): predictions list

    Returns:
        pd.DataFrame: f1-scores dataframe on every class name
    """
    classification_report_dict = classification_report(y_labels, y_true,
→output_dict=True)
    # Create empty dictionary
    class_f1_scores = {}
    # Loop through classification report dictionary items
    for k, v in classification_report_dict.items():
        if k == "accuracy": # stop once we get to accuracy key
            break
        else:
            # Add class names and f1-scores to new dictionary
            class_f1_scores[class_names[int(k)]] = v["f1-score"]
    class_f1_scores

    # Trun f1-scores into dataframe for visualization
    f1_scores = pd.DataFrame({"class_names": list(class_f1_scores.keys()),
                             "f1-score": list(class_f1_scores.values())}).
→sort_values("f1-score", ascending=False)
    return f1_scores


f1_scores = get_f1_score_on_every_class_name(y_labels=y_labels,
→y_true=pred_classes, class_names=class_names)
f1_scores
```

```
[37]: class_names  f1-score
1      forest    0.986272
5      street    0.914697
4       sea     0.914230
0  buildings    0.905077
2    glacier    0.848000
3    mountain    0.844575
```

```
[38]: import matplotlib.pyplot as plt
import pandas as pd

def plot_f1_scores_on_every_class_name(f1_scores, figsize = (10, 10)):
    fig, ax = plt.subplots(figsize=figsize)
    scores = ax.barh(range(len(f1_scores)), f1_scores["f1-score"].values)
    #ax.bar_label(scores, label_type='center', c="white")
    ax.set_yticks(range(len(f1_scores)))
    ax.set_yticklabels(f1_scores["class_names"])
    ax.set_xlabel("F1-score")
    ax.set_title("F1-scores for 101 Different Food Classes")
    ax.invert_yaxis(); # reverse the order of our plot

[39]: plot_f1_scores_on_every_class_name(f1_scores=f1_scores, figsize=(15, 15))
```



```
[40]: # Create a function to load and prepare images
def load_and_prep_image(filename, img_shape=150, scale=True):
    """
    Reads in an image from filename, turns it into a tensor and reshapes into
    specified shape (img_shape, img_shape, color_channels=3).

    Args:
        filename (str): path to target image
        image_shape (int): height/width dimension of target image size
        scale (bool): scale pixel values from 0-255 to 0-1 or not

    Returns:
        Image tensor of shape (img_shape, img_shape, 3)
    """
    # Read in the image
    img = tf.io.read_file(filename)

    # Decode image into tensor
    img = tf.io.decode_image(img, channels=3)

    # Resize the image
    # img = tf.image.resize(img, [img_shape, img_shape])
    img = tf.image.resize(img, list(IMAGE_SIZE))

    # Scale? Yes/no
    if scale:
        # rescale the image (get all values between 0 and 1)
        return img/255.
    else:
        return img # don't need to rescale images for EfficientNet models in
→TensorFlow
```

```
[41]: # Make preds on a series of random images
import os
import random

PRED_DIR: str = "/content/seg_pred/seg_pred"
plt.figure(figsize=(17, 10))
for i in range(3):
    # Choose random image(s) from random class(es)
    filename = random.choice(os.listdir(PRED_DIR))
    filepath = PRED_DIR + "/" + filename
    print(filepath)
    # Load the image and make predictions
    img = load_and_prep_image(filepath, scale=False)
    img_expanded = tf.expand_dims(img, axis=0)
    pred_prob = model_1.predict(img_expanded)
    pred_class = class_names[pred_prob.argmax()]

    plt.subplot(1, 3, i+1)
    plt.imshow(img/225.)
    plt.title(f"pred: {pred_class}, prob: {pred_prob.max():.2f}")
    plt.axis(False);
```

/content/seg\_pred/seg\_pred/5802.jpg

1/1 [=====] - 1s 1s/step

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

/content/seg\_pred/seg\_pred/20760.jpg

1/1 [=====] - 0s 20ms/step

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

/content/seg\_pred/seg\_pred/19330.jpg

1/1 [=====] - 0s 19ms/step

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

